

# OPEN SOURCE MONITORING

Using TIG (Telegraf + InfluxDB + Grafana)

# MONITORING BACKGROUND

These are the are 3 'legs' of the Monitoring Stool:

- Metrics
- Traces
- Logs

Each has its own history, toolings, and uses

Some platforms bring all of these together (known as *Unified Observability*)

It is possible to derive some from the other (e.g. determining metrics from Logs or Traces, searching Traces like Logs)

# Logs

We all know what logs are - the granddaddy of the group

Generally a time-stamped 'message' written to a local file from an application, script or process

There can be varying structure to the log line, and/or the message itself can be structured

Using Logs in *Monitoring* involves shipping, transforming, storing, indexing, and analyzing these log messages

A well-known open-source platform for this is: ELK

- Elasticsearch+Logstash+Kibana from [elastic.co](https://www.elastic.co)

# Traces / Application Tracing

Rise in popularity of *tracing* is tied to the use of services and microservices in building applications

Transactions *through* an application often touch multiple services

Tracing is a strategy to instrument all the services with the ability to emit *spans* for each of these transaction components in and out the service, and collectively multiple spans make up an application *trace*

Best known tooling in this area is Jaeger, but most practitioners are moving towards an accepted standard known as OTel or OpenTelemetry.

# Metrics

We have all interacted with metrics, in one fashion or another.

For example - `top` - emits *metrics* - "numbers that have names"

That is what a metric is - a number. *"...that by which anything is measured."*

To organize the metric, we associate a name with it. That results in the generally accepted format of `{"name" : "load_5", "value" : 0.223}` or even `("load_5",0.223)`

Imagine a system that would just send `("load_5",0.223)` over UDP.

That is the essence of 'statsd' developed by an Etsy engineer in 2011.

Hey! No timestamp! That seems crazy.

# Metrics

- in the original statsd, the timestamp was added by the receiver (keeping the concept of the sender CRAZY simple)
- the receiver could even do aggregations - sending the aggregations 'upstream' every minute and sender could send just whenever (flexible)
- even early on, emitting and receiving metrics was robust and open to innovation
- in general though, a standard 'metric' became something like:

env=prod,region=us-east,name=sys-01 cpu\_load\_5=0.2445 1681179840

0 or more *tags*, followed by a metric name, value, and lastly a timestamp.

# Metrics

what the heck is a *tag*?

- like a category, or a grouping - something that doesn't really change.
- the 'by' in *slice and dice by...*
- it's how you'd answer "what is the average `cpu_load_5` for all my prod servers?"

# TIG (Telegraf - InfluxDB - Grafana)

This is an open-source 'stack' that executes all the important roles in a full monitoring system:

- get lots of great metrics from all sorts of systems
- store them in a tool that is designed for time-series data and has built-in query functions for stats and grouping
- display these metrics in useful and creative ways, using those queries and basic graphing functions
- alert on certain specified conditions of one or more metrics over time



# TIG

## Telegraf:

- compiled golang agent with 'plug-in' pipelined architecture for retrieving, transforming, filtering, and output of metrics

## Influxdb:

- time-series database with advanced query and analysis functions

## Grafana:

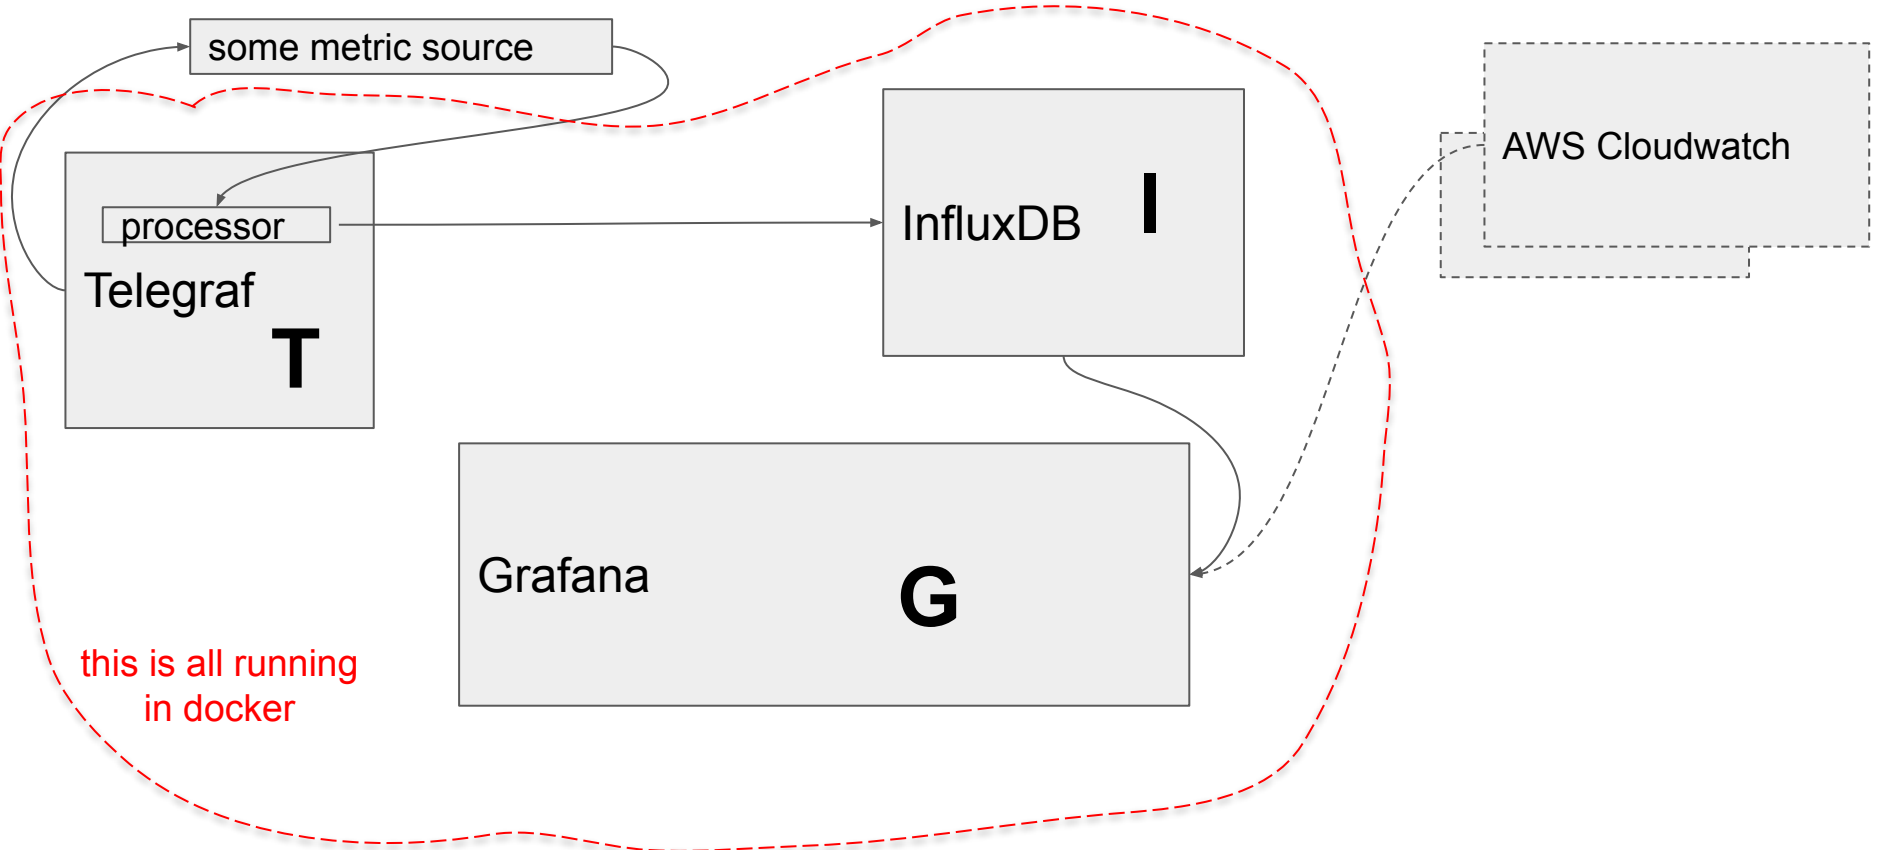
- dashboard and charting tool that also has alerting capabilities

# TIG *stack*

When you get all those working *together*

inspired by this [YouTube](#) where it is all built using *docker-compose*

# TIG architecture



# docker-compose.yml

```
version: "3"

services:
  influxdb:
    image: influxdb:2.1.1
    volumes:
      - influxdb-storage:/var/lib/influxdb2:rw
    env_file:
      - .env
    entrypoint: ["/entrypoint.sh"]
    restart: on-failure:10
    ports:
      - ${DOCKER_INFLUXDB_INIT_PORT}:8086

  telegraf:
    image: telegraf:1.25.3
    volumes:
      - ${TELEGRAF_CFG_PATH}:/etc/telegraf/telegraf.conf:rw
#   networks:
#     mynetwork:
#       ipv4_address: 192.168.86.11
    env_file:
      - .env
    depends_on:
      - influxdb

  grafana:
    image: grafana/grafana-oss:8.4.3
    volumes:
      - grafana-storage:/var/lib/grafana:rw
    depends_on:
      - influxdb
    ports:
      - ${GRAFANA_PORT}:3000

volumes:
  grafana-storage:
  influxdb-storage:

# networks:
#   mynetwork:
#     driver: bridge
#     ipam:
#       config:
#         - subnet: 192.168.86.0/24
```

# .env file

```
DOCKER_INFLUXDB_INIT_MODE=setup
```

```
## Environment variables used during the setup and operation  
of the stack
```

```
#
```

```
# Primary InfluxDB admin/superuser credentials
```

```
#
```

```
DOCKER_INFLUXDB_INIT_USERNAME=admin
```

```
DOCKER_INFLUXDB_INIT_PASSWORD=password
```

```
DOCKER_INFLUXDB_INIT_ADMIN_TOKEN=862010a1b8040e5f2eedc936543  
58d2b24d42cdf1d83d571e2490a4556030a08
```

```
# Primary InfluxDB organization & bucket definitions
```

```
#
```

```
DOCKER_INFLUXDB_INIT_ORG=kyoss
```

```
DOCKER_INFLUXDB_INIT_BUCKET=telegraf
```

```
# Primary InfluxDB bucket retention period
```

```
#
```

```
# NOTE: Valid units are nanoseconds (ns), microseconds(us),  
milliseconds (ms)
```

```
# seconds (s), minutes (m), hours (h), days (d), and weeks  
(w).
```

```
DOCKER_INFLUXDB_INIT_RETENTION=4d
```

```
# InfluxDB port & hostname definitions
```

```
#
```

```
DOCKER_INFLUXDB_INIT_PORT=8086
```

```
DOCKER_INFLUXDB_INIT_HOST=influxdb
```

```
# Telegraf configuration file
```

```
#
```

```
# Will be mounted to container and used as telegraf  
configuration
```

```
TELEGRAF_CFG_PATH=./telegraf/telegraf.conf
```

```
# Grafana port definition
```

```
GRAFANA_PORT=3000
```

# Telegraf

- huge library of plug-ins
- input plug-ins
-

## output plug-ins

this one is configured to send all metrics to influxdb

```
997
998 # # Configuration for sending metrics to InfluxDB
999 [[outputs.influxdb_v2]]
1000     ## The URLs of the InfluxDB cluster nodes.
1001     ##
1002     ## Multiple URLs can be specified for a single cluster, only ONE of the
1003     ## urls will be written to each interval.
1004     ##   ex: urls = ["https://us-west-2-1.aws.cloud2.influxdata.com"]
1005     urls = ["http://${DOCKER_INFLUXDB_INIT_HOST}:8086"]
1006
1007     ## Token for authentication.
1008     token = "${DOCKER_INFLUXDB_INIT_ADMIN_TOKEN}"
1009
1010     ## Organization is the name of the organization you wish to write to; must exist.
1011     organization = "${DOCKER_INFLUXDB_INIT_ORG}"
1012
1013     ## Destination bucket to write into.
1014     bucket = "${DOCKER_INFLUXDB_INIT_BUCKET}"
1015
```